

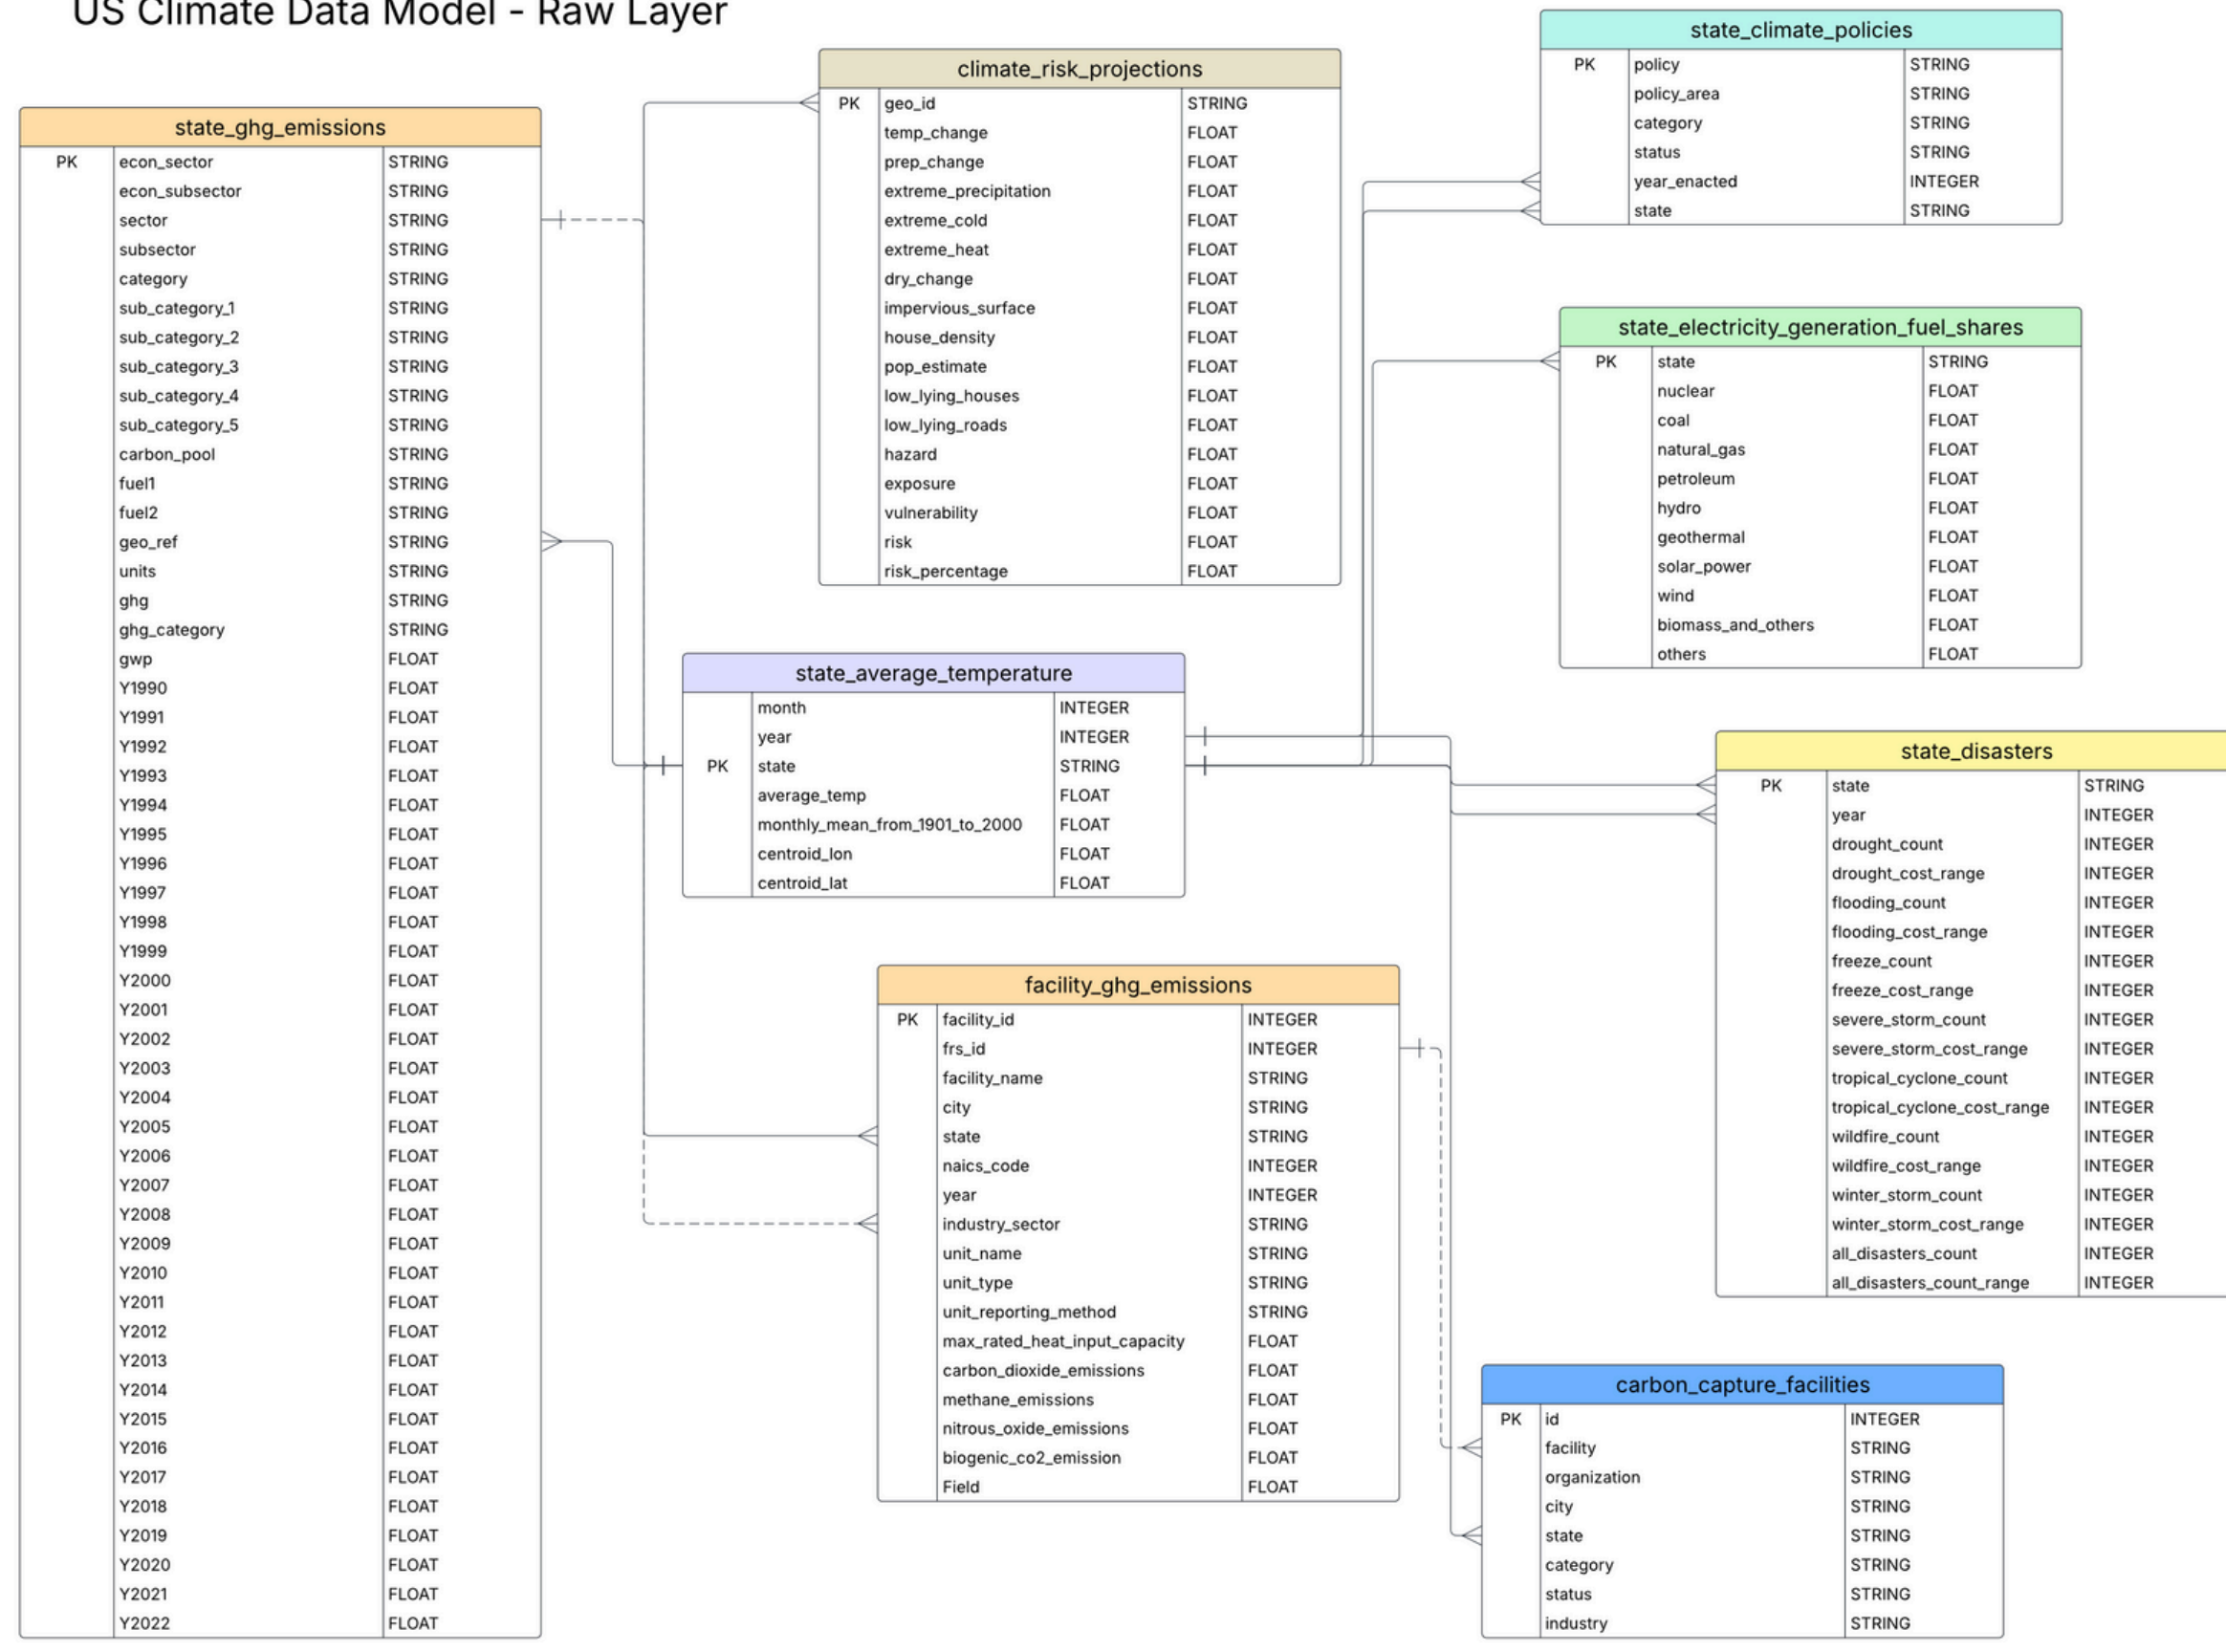


U.S. CLIMATE DATA WAREHOUSE

Erica Zhao, Kiara Foght

RAW DATA

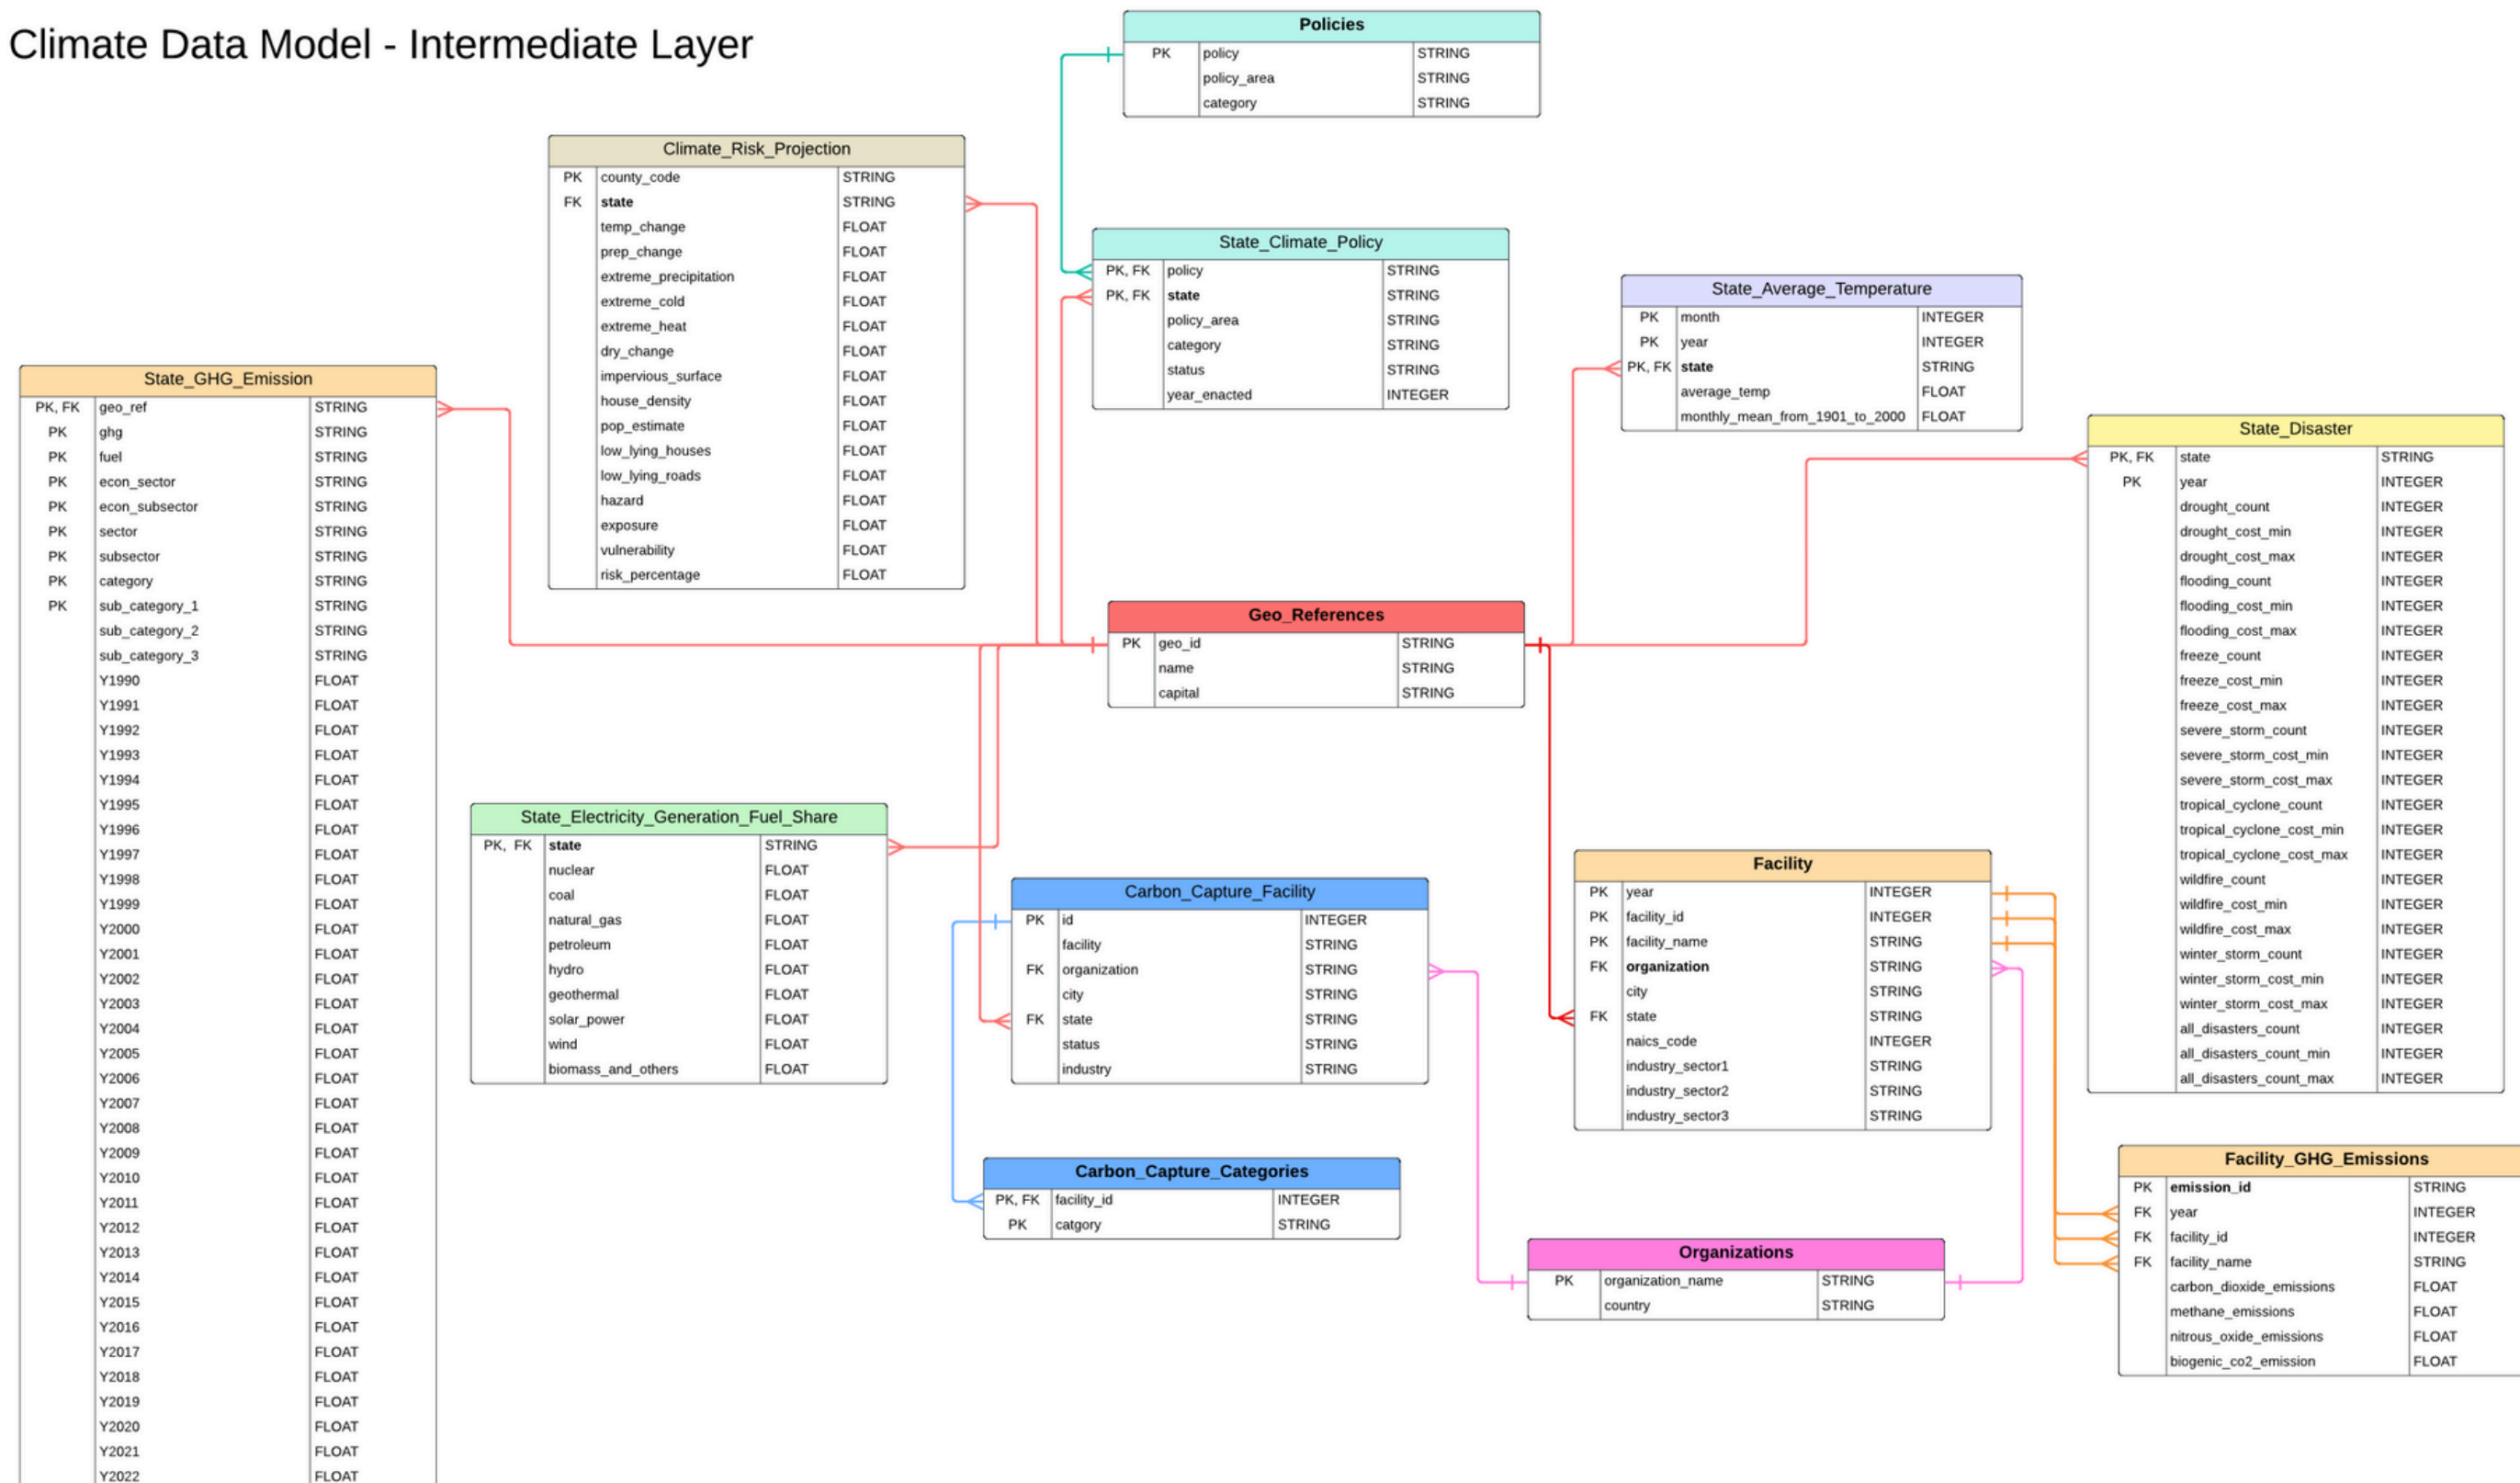
US Climate Data Model - Raw Layer



- EPA
- Kaggle
- Climate XChange
- NOAA
- NEI
- BBI International
- NASA

INT ERD

US Climate Data Model - Intermediate Layer



DBT CHALLENGES

old_State_GHG_Emission

```
21:33:03 11 of 11 START test unique_combination_of_columns_old_State_GHG_Emission_geo_ref__ghg__fuel__sector__subsector__econ_sector__category_1 [RUN]
21:33:03 8 of 11 FAIL 5366 not_null_old_State_GHG_Emission_sub_category_1 ..... [FAIL 5366 in 0.99s]
21:33:04 11 of 11 PASS unique_combination_of_columns_old_State_GHG_Emission_geo_ref__ghg__fuel__sector__subsector__econ_sector__category_1 [PASS in 1.09s]
21:33:04 10 of 11 PASS relationships_old_State_GHG_Emission_geo_ref__geo_id__ref_GeoReferences_ [PASS in 1.17s]
21:33:04
21:33:04 Finished running 11 data tests in 0 hours 0 minutes and 4.35 seconds (4.35s).
21:33:04
21:33:04 Completed with 4 errors, 0 partial successes, and 0 warnings:
21:33:04
21:33:04 Failure in test not_null_old_State_GHG_Emission_econ_subsector (models/int/schema.yml)
21:33:04   Got 2200 results, configured to fail if != 0
21:33:04
21:33:04   compiled code at target/compiled/us_climate/models/int/schema.yml/not_null_old_State_GHG_Emission_econ_subsector.sql
21:33:04
21:33:04 Failure in test not_null_old_State_GHG_Emission_category (models/int/schema.yml)
21:33:04   Got 28 results, configured to fail if != 0
21:33:04
21:33:04   compiled code at target/compiled/us_climate/models/int/schema.yml/not_null_old_State_GHG_Emission_category.sql
21:33:04
21:33:04 Failure in test not_null_old_State_GHG_Emission_fuel (models/int/schema.yml)
21:33:04   Got 17049 results, configured to fail if != 0
21:33:04
21:33:04   compiled code at target/compiled/us_climate/models/int/schema.yml/not_null_old_State_GHG_Emission_fuel.sql
21:33:04
21:33:04 Failure in test not_null_old_State_GHG_Emission_sub_category_1 (models/int/schema.yml)
21:33:04   Got 5366 results, Open file in editor \(cmd + click\)
21:33:04
21:33:04   compiled code at target/compiled/us_climate/models/int/schema.yml/not_null_old_State_GHG_Emission_sub_category_1.sql
21:33:04
21:33:04 Done. PASS=7 WARN=0 ERROR=4 SKIP=0 TOTAL=11
```


DBT FIX

State_GHG_Emission

```
-- In the original State_GHG_Emission table the various fields set to be primary keys had nulls so tests failed
-- To fix this :
-- 1. Unpivot table so the multiple years are not each a separate field
-- 2. Summed emission values when all field values were the same
-- 3. Created emission_id to use as PK that is a string of all the separate fields that you have been PK's
```

```
with int_tmp_state_ghg_unpivot as (
  select
    geo_ref,
    cast(substring(year, 2) as int64) as year,
    ghg,
    fuel,
    sector,
    subsector,
    econ_sector,
    econ_subsector,
    category,
    sub_category_1,
    sub_category_2,
    sub_category_3,
    emission,
    _data_source,
    _load_time
  from {{ ref('state_ghg_emissions') }}
  unpivot(
    emission for year in (
      Y1990, Y1991, Y1992, Y1993, Y1994, Y1995,
      Y1996, Y1997, Y1998, Y1999, Y2000, Y2001,
      Y2002, Y2003, Y2004, Y2005, Y2006, Y2007,
      Y2008, Y2009, Y2010, Y2011, Y2012, Y2013,
      Y2014, Y2015, Y2016, Y2017, Y2018, Y2019,
      Y2020, Y2021, Y2022
    )
  )
),

int_tmp_state_ghg_aggregated as (
  select geo_ref, year, ghg, fuel, sector, subsector, econ_sector, econ_subsector,
    category, sub_category_1, sub_category_2, sub_category_3, sum(emission) as emission,
    _data_source, _load_time
  from int_tmp_state_ghg_unpivot
  group by geo_ref, year, ghg, fuel, sector, subsector, econ_sector, econ_subsector,
    category, sub_category_1, sub_category_2, sub_category_3, _data_source, _load_time
),
```

```
int_State_GHG_Emission as (
  select
    array_to_string([
      geo_ref,
      cast(year as string),
      ghg,
      fuel,
      sector,
      subsector,
      econ_sector,
      econ_subsector,
      category,
      sub_category_1,
      sub_category_2,
      sub_category_3
    ], ', ') as emission_id,
    geo_ref,
    year,
    ghg,
    fuel,
    sector,
    subsector,
    econ_sector,
    econ_subsector,
    category,
    sub_category_1,
    sub_category_2,
    sub_category_3,
    emission,
    _data_source,
    _load_time,
  from int_tmp_state_ghg_aggregated
)

select *
from int_State_GHG_Emission
```

DBT FIX

State_GHG_Emission

```
(dbt-env) kmfoght@dbt01:~/us_climate$ dbt test --select State_GHG_Emission.sql
21:51:48 Running with dbt=1.9.3
21:51:49 Registered adapter: bigquery=1.9.1
21:51:50 Found 44 models, 57 data tests, 9 sources, 494 macros
21:51:50
21:51:50 Concurrency: 3 threads (target='dev')
21:51:50
21:51:50 1 of 3 START test not_null_State_GHG_Emission_emission_id ..... [RUN]
21:51:50 2 of 3 START test relationships_State_GHG_Emission_geo_ref__geo_id__ref_Geo_References_ [RUN]
21:51:50 3 of 3 START test unique_State_GHG_Emission_emission_id ..... [RUN]
21:51:51 1 of 3 PASS not_null_State_GHG_Emission_emission_id ..... [PASS in 1.15s]
21:51:52 2 of 3 PASS relationships_State_GHG_Emission_geo_ref__geo_id__ref_Geo_References_ [PASS in 1.31s]
21:51:52 3 of 3 PASS unique_State_GHG_Emission_emission_id ..... [PASS in 2.20s]
21:51:53
21:51:53 Finished running 3 data tests in 0 hours 0 minutes and 2.54 seconds (2.54s).
21:51:53
21:51:53 Completed successfully
21:51:53
21:51:53 Done. PASS=3 WARN=0 ERROR=0 SKIP=0 TOTAL=3
```

DBT

Final State_GHG_Emission

State_GHG_Emission		
PK	emission_id	STRING
FK	geo_ref	STRING
	year	INTEGER
	ghg	STRING
	fuel	STRING
	sector	STRING
	subsector	STRING
	econ_sector	STRING
	econ_subsector	STRING
	category	STRING
	sub_category_1	STRING
	sub_category_2	STRING
	sub_category_3	STRING
	emission	FLOAT

ORGANIZATION TABLE

Use LLM to Identify Organizations

Normalize Organization Names

Create the Final Organization Table

Validate and Finalize the Organization Table

LLM

Used a large language model to semantically evaluate organization names and return the most representative or distinguish between unrelated ones.

```
prompt = f"""
Normalize the organization name and determine the country this organization belongs to.
For example, X and X LLC should be the same company X, Z glass and Z group should be the same company.
Given the organization: "{org}"

Return JSON in this format:
{{
  "original_name": "{org}",
  "standardized_name": "<Standardized Organization Name>",
}}

No extra text or explanation. Only return valid JSON.
"""
```

VS

Fuzzy Matching

Converted names into embeddings and grouped similar ones using cosine distance, then reviewed borderline cases with LLM.

FUZZY MATCHING

01

Generate Text
Embeddings

02

Perform
Nearest
Neighbor
Search

03

Apply a
Distance
Threshold

```
[ ] %%bigquery

CREATE OR REPLACE TABLE us_climate_fin.ghg_org_embeddings AS (

WITH org_content AS (
  SELECT
    organization_name,
    organization_name AS content
  FROM
    us_climate_fin.ghg_org_names_raw
)

SELECT
  organization_name,
  content,
  ml_generate_embedding_result AS embedding
FROM
  ML.GENERATE_EMBEDDING(
    MODEL us_climate_fin.embedding_model,
    (
      SELECT organization_name, content
      FROM org_content
      WHERE content IS NOT NULL
    ),
    STRUCT('CLUSTERING' AS task_type)
  )
)
```

```
[ ] %%bigquery
select *
from us_climate_fin.ghg_org_nearest_neighbors
where distance <= 0.075
order by distance
```

```
[ ] %%bigquery

CREATE OR REPLACE TABLE us_climate_fin.ghg_org_nearest_neighbors AS

SELECT
  query.organization_name AS organization_name,
  base.organization_name AS nearest_neighbor,
  distance
FROM
  VECTOR_SEARCH(
    TABLE us_climate_fin.ghg_org_embeddings,
    'embedding',
    TABLE us_climate_fin.ghg_org_embeddings,
    'embedding',
    TOP_K => 2,
    DISTANCE_TYPE => 'COSINE'
  )
WHERE
  query.organization_name != base.organization_name
ORDER BY
  distance
```

FUZZY MATCHING

```
import pandas as pd
import pandas_gbq
from google.cloud import bigquery
input_table = "us_climate_fin.ghg_org_nearest_neighbors"
output_table = "us_climate_fin.ghg_org_clusters"
base_query = f"""
    SELECT organization_name, nearest_neighbor
    FROM `{input_table}`
    WHERE distance <= 0.075
"""
bq_client = bigquery.Client()
rows = bq_client.query(base_query).result()

cluster_id = 0
output_clusters = []
unique_ids = set()
for row in rows:
    id1 = row["organization_name"]
    id2 = row["nearest_neighbor"]
    if id1 not in unique_ids and id2 not in unique_ids:
        cluster_id += 1
        output_clusters.append((id1, cluster_id))
        output_clusters.append((id2, cluster_id))
        unique_ids.add(id1)
        unique_ids.add(id2)
        print(f"Assigned {id1} and {id2} to cluster {cluster_id}")
df = pd.DataFrame(output_clusters, columns=["organization_name", "cluster_id"])
pandas_gbq.to_gbq(df, output_table, project_id="kiaraerica", if_exists="replace")
```

04

Assign Cluster
IDs

05

Choose
Standardized
Name

```
prompt = """Please verify each organization name.
If they are under the same parent organization, return the closest to the parent
organization name in json using the schema {"organization_name" : string}.
For example, {"organization_name": "EPL Oil & Gas, Inc."}
If they refer to different parent organizations, return both organization names
in json using the schema
[{"organization_name": "<name1>"}, {"organization_name": "<name2>"}]
For example,
[{"organization_name": "University of Utah"},
{"organization_name": "Utah State University"}]
Do not include an explanation with your answer.
"""
```

organization_name	nearest_neighbor	distance
EPL Oil and Gas, Inc.	EPL Oil & Gas, Inc.	0.001325127...
EPL Oil & Gas, Inc.	EPL Oil and Gas, Inc.	0.001325127...


organization_name	nearest_neighbor	distance
University of Utah	Utah State University	0.050155917161...
Utah State University	University of Utah	0.050155917161...

COMPARISION

This method captured duplicates missed by traditional string matching, especially those with abbreviations or inconsistent formatting.

This let us group similar names before invoking the LLM.

However, to ensure reasonable recall, we had to use a relatively large distance threshold, which introduced some false positives and still missed a few known duplicates.



Fine-tune the trade-off & Link the normalized organization names back to their corresponding facility IDs

**THANK
YOU**

